

CS-202 Week 1

Notes

Lecture 1

Admin & Goals

Com Sys is a new course replacing the following:

- Introduction to Computer Systems, CS-323
- Computer Networks, COM-208
- Programmation Orienté Systems, CS-207
- Projet de programmation orienté Systems, CS-212

Course Aims:

- Understanding timeless principles of os and networking
- Understand the difference between application-oriented and system-oriented thinking
- Separate the ends, understand middle-boxes and routers
- Build hands on experience

Requirements:

- Linux os or VM

Structure:

- Main Lectures:
 - 4 hours of non recorded / streamed lectures
 - Slides available before-hand
 - Summary videos available
 - Weekly graded quiz to check understanding
- C Bootcamp (first 3 weeks):
 - Flipped teaching
 - 3 weeks theory -> practice
 - 2 person graded programming project starting week 4

The role of the operating system

[Operating System \(OS\)](#) is a special type of program, it should never stop and never fail. It has a trust value.

Hardware Foundations of Computer Systems

The von Neumann Architecture:

CPU:

- Arithmetic/Logic Unit: Does operations

- Control Unit: Controls what happens next
 - has at least one register, the instruction pointer (IP):
 - interprets instructions at address IP
- Programs are stored in memory

Protection and memory management units

Key enhancements to the von neumann:

CPU privileges:

- user mode
- kernel mode (special access)

Virtual memory with a [Memory Management Unit \(MMU\)](#):

- MMU converts virtual addresses to physical addresses

CPU operates with virtual addresses only:

- Instruction pointer is a virtual address
- All pointers are virtual addresses

What is an OS

An OS is a software that **interfaces** between the physical hardware resources and the one or many apps running on the machine. Implements **abstraction** that is easier to program than the raw hardware.

OS wears three main hats:

- Referee: manages protection, isolation and sharing of resources
- Illusionist: Provide abstraction of physical resources
- Glue: Provide a set of common resources

The OS as a referee

Full isolation:

- Isolate programs from each other and from OS
- Resource sharing:
- Choosing which program to execute
 - Splitting physical resources for programs
- Communication:
- Primitives that allow communication among programs

The OS as a glue

Makes sharing easier

Maximise reuse:

- avoid re-implementation
- decouple hardware and software development

Lecture 2

2.1 - The process abstraction

Focus of lecture

the need of a process abstraction

what is a process

how os creates a process

Abstractions provide an **interface** to application programmers that separates **policy** - what the interface commits to accomplishing - from **mechanism** - how the interface is implemented

Executable file

Contains:

- Executable code: CPU instructions
- Data: Information manipulated by these instructions

What constitutes a process ?


- A unique identifier: Process ID (PID)
- Memory image (contains 4 *segments):
 - Code and data (static)
 - Stack and heap (dynamic)
- CPU context: registers
 - Program counter, current operands, stack pointer
- File descriptors: pointers to open files and descriptors

Running a program -> creating a process

a process is an instance of an executable

The register file of a process can be in memory or in CPU but not in both at the same time.

What is in the process memory



Stack

- Temporary data like function parameters, local vars and return addresses
- Grows from higher to lower addresses
- How does Stack Pointer (SP) know when it has reached heap ?
 - OS protects SP from pointing to heap
 - Stack overflow: when stack grows beyond allocated memory due to too many function calls

Heap

- Used for dynamic memory allocation during program runtime

Data

- Statically allocate global variables and program data structures

Text

- Read-only
- Contains code and constants
- These are program executable machine instructions

Stack is necessary because of recursion

Heap is where we allocate objects that can live outside a function

Null pointer:

How to OS creates a [process](#)

- Loading: OS loads the static code and data into memory
- Memory Allocation: allocate process memory regions (heap and stack)
- Initialization: Initialize tasks related to IO (setting up STDIN, STDOUT, STDERR)

- Ready: OS sets up the stage for running the process by transferring CPU control to beginning of program entry point

Sharing of resources: Two approaches

Time sharing

Running one task at a time and quickly switching between resources

Space sharing

Each task gets a portion of available space


OS time shares CPU among multiple processes

2.2 - Process state execution diagram

State transition

3 states:

- Running: process is running on CPU
- Ready: ready to run but OS has not chosen to run it yet
- Blocked: Not ready to run until some event occurs



OS mechanism for process management

Process control block

Information for each process is stored in a process control block (slide 27)

Process API

pid 1: init

fork()

- Process calling fork is parent process
- creates a child process which is duplicate of parent process
- returns child pid
- (continue)

exec()

- Often used not long after fork
- Loads a new program in the context of an already running process (child), replacing the previous executable program
- replaces the memory (address space) - code data heap and stack are replaced
- No new PID
- STDIN, STDOUT and STDERR are kept to allow parent to redirect/rewire child's output

wait()

- blocking call
- wait for child to finish executing

exit()

- exits process