

# CS-202 Week 2

## Notes

### Lecture 1:

### Isolation and Protection

#### Focus:

- Limited direct execution
- Protecting processes from each other

#### Protection

- Protect applications from other application's code
    - ensures reliability, security and privacy
  - Protect the os from applications
  - Protect applications from unfavorable resource utilization
- OS is fully trusted, whilst other applications are untrusted

OS allows processes to run with **restricted rights**

Enforces mechanisms that allow processes to run with **potentially dangerous operations disabled**

OS enforces time-sharing even when the process does not cooperate (example infinite loop)

### CPU virtualization: efficient process execution

**Goal:** give each process the illusion of an exclusive CPU access

**Reality:** CPU is shared among all processes

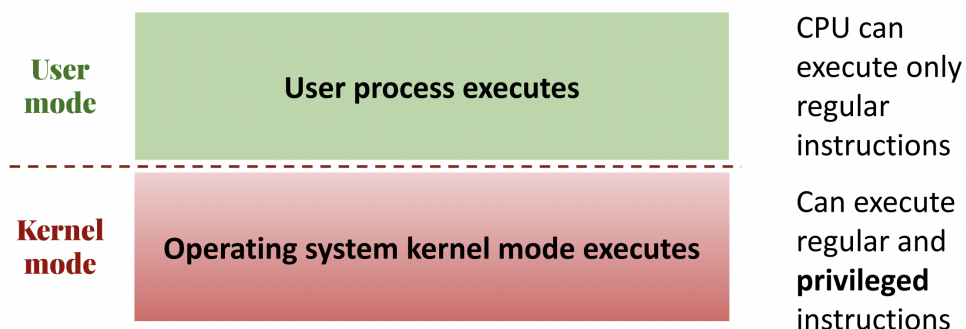
CPU uses time sharing

### Limited Direct Execution

Enables programs to execute as fast as possible with some restrictions.

(**Direct execution:** run programs directly on CPU without any restrictions.)

**Limited direct execution with dual mode:**



### Privileged instructions:

Examples:

- Access I/O devices
- Change privilege levels
- Enable or disable interrupts
- Change the MMU register that controls page table (mov %cr3 on x86-64)
- ...

When CPU tries to execute privileged instruction from user mode:

- The instruction does not execute
- The instruction traps (#General protection fault on x86)
- The operating system takes control

## System calls

Processes can request OS services through the **system call API**


System calls transfer execution to the OS, meanwhile execution of the process is suspended

Processes execute special **trap instructions** to execute system calls (slide 28):

- CPU jumps into kernel mode and changes privilege level at the same time
- Now privileged operations can be performed

**Trap** is a signal that instructs OS to perform some functionality immediately

Slide 32 (important)



## Traps

Traps also referred to as exceptions.

Exceptions produced by CPU while executing instructions.

Exceptions are synchronous: invoked by CPU only after termination the invocation of an instruction.

OS configures hardware at boot time (slide 36)

Requesting OS services using system call numbers (slide 37)

## Interrupts

Hardware provides signaling mechanism for the OS to regain control.

Interrupts are asynchronous signals to the CPU that some external event has occurred and requires attention.

Example: Timer interrupt, IO interrupt, inter-processor interrupt

OS configures the **timer interrupt** to regain control

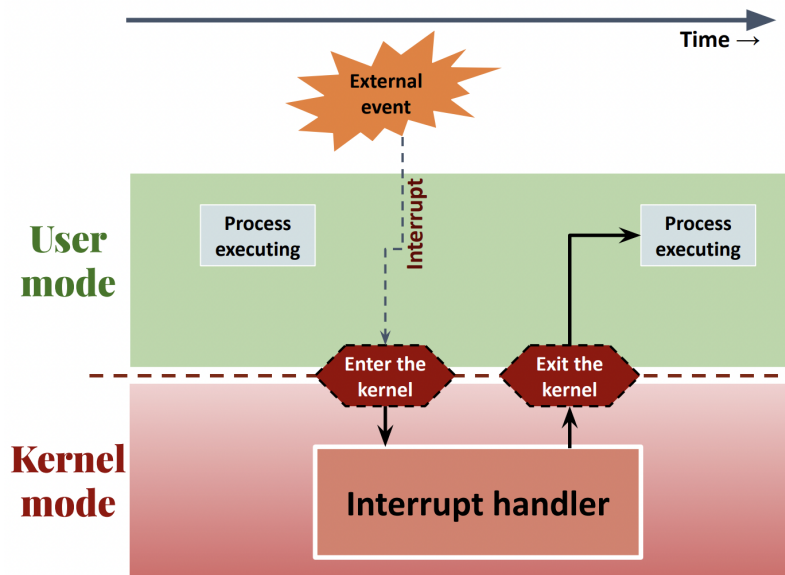
Interrupt is fired every certain milliseconds (every 10 ms in Linux kernel)

With every interrupt, CPU switches from user to kernel mode (context switch) so OS can choose which process to execute

## Handling

When there is an interrupt:

- The CPU detects it
- Suspends the user process and switches to kernel mode
- Executes the appropriate interrupt handling code
- Restores the user process



## Concurrent interrupts

Hardware provides instructions to temporarily delay the delivery of an interrupt (**disable interrupt**), and enable it again when safe

Interrupts are disabled when an interrupt handler is running

## Lecture 2: The os scheduler

### Scheduling mechanisms

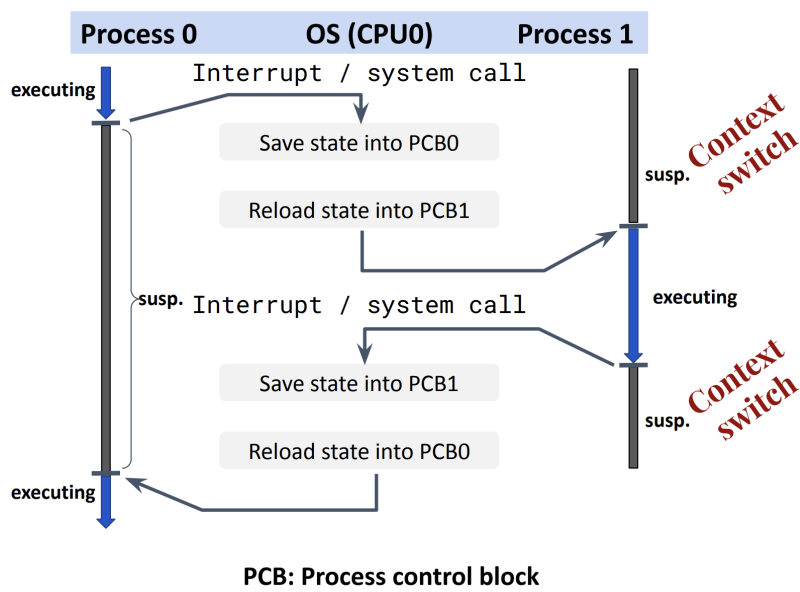
#### Context switching

- OS in kernel mode and cannot return to same process
- OS does not want to run same process
- -> OS performs a **context switch**

A mechanism that allows the OS to store the current process state and switch to another previously stored context

Context of process is represented in the process control block (PCB)

OS maintains PCB for each process



Preemption: OS sets a timer before scheduling a process. Hardware generates interrupt after timer expires

## Scheduling metrics

### Utilization

What fraction of the time is the CPU executing a job

**Goal:** Maximize CPU utilization as much as possible

### Turnaround time

Total time from a job arrival to job completion

**Goal:** Minimize the total time as much as possible

### Response time

How long long a job takes to be scheduled

**Goal:** Minimize total time as much as possible

## The idle process

Low priority process that is scheduled if no other process is ready.

What process to schedule if all processes are blocked ? The idle process (runs in kernel mode).

Idle mode issues special instructions to save power.

## CPU-only Scheduling Policies (not used)

### Policies that optimise turnaround time

#### FIFO (First in first out)

Not a viable strategy. Long jobs delay short jobs (example: slide 29)

Only works is processes have the same amount of execution time.

convoy effect (see slides)

#### SJF (Shortest job first)

Schedule shortest job first and run to completion.

Reduces average turnaround time.

Long running jobs cannot be interrupted.

Still subject to convoy effect if longer process arrives before others.

## **STCF (Shortest Time to Completion First)**

FIFO and SJF are non-preemptive: only switch to other job when current is finished

STCF extends the SJF by adding preemption

Any time a new job enters system, STCF Scheduler determines which of remaining job has shortest completion time.

STCF then schedules the shortest job first.

Response is not great...

## **Policies that optimise response time**

### **RR (Round Robin)**

Runs jobs for a fixed interval, performing a context switch after said time if other jobs are available

RR policy is fair, evenly divides CPU time among ready processes on a small time scale.

Higher average turnover but lower response time.

## **CPU & I/O scheduling**

### **IO awareness**

Almost every program performs IO of some kind

IO is usually slow: Read/write requests can take milliseconds to complete

Schedulers need to incorporate CPU and IO by blocking processes.

## **MLFQ (Multi-level Feedback Queue)**

### **Goal: Supporting general purpose scheduling**

**First:** MLFQ optimises turnaround time:

- important for batch processes

**Then:** minimizes response time for better interactivity

**Insight:** Use past behavior as predictor for future behavior

### **Approach: Multiple levels of round robin**

Each level has higher priority and preempts lower level

Process at a higher level will always be scheduled first

High levels have short time slices, lower levels run longer

Rule 1: if  $\text{priority}(A) > \text{priority}(B)$  then A runs

Rule 2: if  $\text{priority}(A) == \text{priority}(B)$  then A & B run in RR

Rule 3: Processes start at top priority

Rule 4: If a process uses its total time slice, the scheduler lowers its priority

Rule 5: Periodically moves all jobs to topmost queue (**priority boosting**)

**priority boosting:** jobs moved up to highest priority queue

Boosting avoids starvation on long running jobs

